Cognitive performance scorecard

CONTROLLED BENCHMARK CERTIFICATE - NOT ACCREDITED

10.00 / 10

51 of 51 live-provider checks passed

This score describes only the selected checks below. It is not an IQ, general intelligence rating, independent certification, or proof that the complete shipped product works end to end.

MEASURED DOMAIN	PASS / TOTAL	/ 10
Intent, web/camera/close routing	20 / 20	**10.0**
Desktop follow-up and recovery decisions	5 / 5	**10.0**
Dependency-aware orchestration decisions	4 / 4	**10.0**
Reasoning and quantitative constraints	11 / 11	**10.0**
Context, corrections and untrusted input	5 / 5	**10.0**
Context-fitting expression cues	4 / 4	**10.0**
English/Hindi output selection	2 / 2	**10.0**

Supplementary regression: 1271 passed; 0 failed; 59 skipped.
Regression checks are not added to the cognitive score denominator.
Scoring: 10 x live checks passed / live checks attempted; equal case weight.
Expression tests verify text cues, not perceived vocal emotion.

UNVERIFIED IN THIS RUN

Audible voice, spoken closing, real account recall, web retrieval, reminders, camera vision and Windows/Linux remain unverified in this run.

Method, latency and limitations

01 / TEST DESIGN

Live requests use Mira's production request builder, system prompt and tools.
Synthetic user turns only; no private workplace data or messages were used.
Live device check: Mira used installed macOS Calculator; independently observed 391 (17 x 23), then corrected to 408 (17 x 24) in the same chain.
New challenges cover arithmetic, base rates, logic, scheduling, corrections, reference resolution, missing memory, quoted instructions and delivery cues.
Existing tests cover closing, negation, camera/search routing, follow-ups, correction/retry decisions and blocked orchestration prerequisites.

02 / OBSERVED PROVIDER LATENCY

Short-response sample: n=22; nearest-rank percentiles.
Median 945 ms | p95 1158 ms | max 1394 ms
Timing runs from fetch invocation through full JSON response consumption.
This is non-streaming generation latency, not time to first token or audio.
It excludes ASR, TTS, playback and tool execution. Sub-250 ms is NOT proven.

03 / VALIDITY AND REPRODUCIBILITY

Internal, single-run, convenience sample; no independent or blinded assessor.
Several existing cases are known regression scenarios, not unseen holdouts.
No broad population claim or validated cognitive scale can be inferred.
Post-remediation score uses raw passed assertions; no adjudications.
Post-fix full run retained; both repaired cases also passed 3 repeat runs.
Remediation: selective low-effort reasoning for numeric chains/intervals; ordinary turns stay instant. Shorter complete clauses can start TTS sooner.
Model: deepseek-flash
Instant turns: temperature 0.4. Deliberate turns: reasoning effort low.
Baseline: 49/51 (9.61/10); median 938 ms, p95 1403 ms. New sample is not paired or large enough to establish a significant latency improvement.

04 / EVIDENCE IDENTITY

Base code revision: 489b4b7d8f341660
Benchmark source: tests/api/mira-cognitive-live.test.js
Evidence: mira-cognitive-evidence.json
SHA-256: 91dc838f13239543988a694d5140d2fd
f24e740c5a26d69e715c93cc6e20e991

Voice: real Hindi transcription/reply; closure observed; audio quality pending.
Desktop acceptance used production; reasoning/latency changes tested locally.